

Dynamic Login Node Resource Control and Monitoring with Arbiter 3

Jackson McKay, Kai Forrest, and Paul Fischer



Problem: Shared Interactive Resource Allocation

Challenge: Provide best quality-of-service to interactive users of a shared resource.

- Fair – No single user hogging a node
- Reliable – User experience consistent
- Bursts – Allow for short bursts of high utilization



Problem: Shared Interactive Resource Allocation

Static Limits?

- Poor utilization
- Not bursts

Proportional Limits?

- Unpredictable UX

Solution: Arbiter monitors resource usage in real-time and gently constraints users that violate resource policies.

- Fair – No single user hogging a node ✓
- Reliable – User experience consistent ✓
- Bursts – Allow for short bursts of high utilization ✓




Arbiter2

Features

- Dynamic scaling
- Notifications


Issues

- V2 support
- Synchronization
- Installation

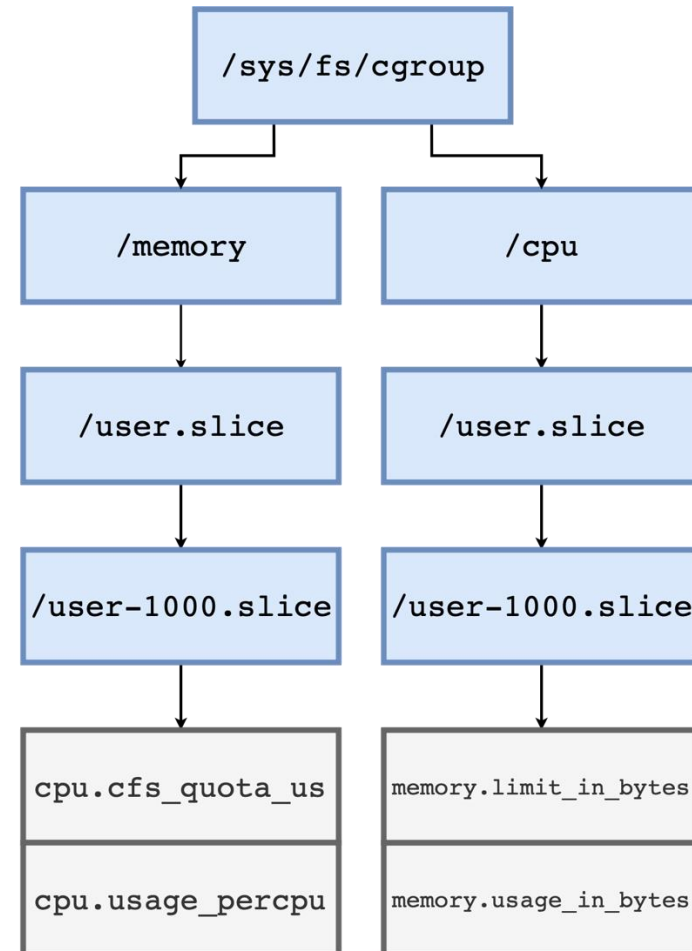


Linux cgroups

Unified (v2)



Legacy (v1)





Improvements

- Deployment and Configuration
 - Web GUI
 - Single source of configuration
 - Interactive nodes run single binary
- Performance
 - Go > Python
 - Offload policy evaluation
- Flexibility
 - Use new or existing Prometheus-based monitoring system
 - Policies are expressed as PromQL queries on arbitrary metrics
 - Policies and systems can be added/removed on the fly



User Evaluation and Enforcement

- Evaluation with Prometheus

- Offload work to TSDB
- Time-series means no proxy 'badness' score

```
rate(systemd_unit_cpu_usage_ns{instance="notchpeak1.chpc.utah.edu"}[15m]) / threshold > 1
```

- Enforcement with cgroup-warden

- Utilizes Systemd properties
- cgroup version agnostic

```
{"unit": "user-1000.slice", "property": {"name": "CPUQuotaPerSecUSec", "value": 1000000000000}}
```



Arbiter | Dashboard

https://arbiter.chpc.utah.edu/dashboard

Dashboard

Base Policy

Usage Policy

Violations

Log Out

Management

Status:

cgroup-agent instances up: 6

Last-Evaluated: 7 seconds ago

Run Arbiter Evaluation Loop:

Evaluate

Clean Violation History:

Before: 11 / 15 / 2024 , 04 : 00 PM

Clean

Apply Property

Host

lonepeak1.chpc.utah.edu

Username

Property

CPU Quota (cores)

Value

0

Apply Limit

Recent Violations

Violation	Time
u1259516@scrubpeak1.chpc.utah.edu - Policy object (3)	8 seconds ago

View Usage

Host: All

Start Time: mm / dd / yyyy , -- : -- --

Step Size: 30

Seconds

View Memory

Username: u1259516

End Time: mm / dd / yyyy , -- : -- --

View CPU

Links and Contact

- Arbiter3
 - <https://github.com/chpc-uofu/arbiter>
- cgroup-warden
 - <https://github.com/chpc-uofu/cgroup-warden>
- Email
 - jay.mckay@utah.edu
 - kai.forrest@utah.edu
 - p.fischer@utah.edu

